

HW2

b09902004 資工二 郭懷元

2.8

```
addi x30, x10, 8 → x30 = A+1
addi x31, x10, 0 → x31 = A
sd x31, 0(x30) → *(x30) = x31 → *(A+1) = A → A[1] = A
ld x30, 0(x30) → x30 = *(x30) → x30 = *(A+1) = A
add x5, x30, x31 → f = A + A
```

```
A[1] = A;
f = A[1] + A;
```

2.9

instruction	opcode	rs1	rd	rs2	imm	funct3	funct7
addi x30, x10, 8	010011	01010	11110		0x008	000	
addi x31, x10, 0	010011	01010	11111		0x000	000	
sd x31, 0(x30)	100011	11110		11111	0x000	011	
ld x30, 0(x30)	000011	11110	11110		0x000	011	
add x5, x30, x31	110011	11110	00110	11111		000	000

2.16

2.16.1

`funct7`, `funct3`, `opcode` : These bit fields might increase in size to accommodate the four times as many instructions.

`rs2`, `rs1`, `rd` : These bit fields should increase from 5 bits to 7 bits for the 128 registers.

2.16.2

`funct3`, `opcode` : These bit fields might increase in size to accommodate the four times as many instructions.

`rs1`, `rd` : These bit fields should increase from 5 bits to 7 bits for the 128 registers.

`imm` : This field doesn't need to change, because neither the number of registers or instructions have to do with `imm`.

2.16.3

Decrease in size: Because there are more registers and more instructions, some old instructions can now be combined into just a single instruction.

Increase in size: Because instructions now takes up more bits, for simple tasks that doesn't use many registers, the extra bits are wasted and take up unnecessary spaces.


Matrix Multiplication Report

Refs:

https://en.wikipedia.org/wiki/Loop_nest_optimization
<https://github.com/flame/how-to-optimize-gemm/wiki>

Result: 8960429 cycles

Strategy



- Dot 2 rows of A and 4 columns of B at a time, so that when a value from A or B is read from memory to register, it can be used multiple times.
- Use pointers and indirect addressing to access A and B .

- I tried memory blocking to keep data in L1 and L2 cache, but somehow it didn't work, and the performance is even worse.

Questions

How many cycles does it take by just doing the naive matrix multiplication?

The naive C implementation with gcc `-O3` optimization takes slightly more than 16M cycles.

How many load and store does it need (roughly) during the whole computation? (Considering the registers it use)

In the naive implementation, C is stored once, A and B are both loaded into register 128 times. So total store is about $2^{7 \times 2} = 2^{14}$ times and total load is about $2 \times 2^7 \times 2^{7 \times 2} = 2^{22}$ times.

With my strategy, C is also stored once, A is loaded $\frac{128}{4} = 32$ times, and B is loaded $\frac{128}{2} = 64$ times. So total store is 2^{14} times, and total load is $2^5 \times 2^{14} + 2^6 \times 2^{14} = 3 \times 2^{19}$ times, $\frac{3}{8}$ of naive implementation.

Is there any way to keep registers being used as much as possible before they're replaced? (Hint: blocking)

Yes. My strategy is intended to keep reusing registers as much as possible. The number of rows and columns we dotted at a time is bounded by how many registers are available.

How many loop controls does it need (roughly) during the whole computation?

Both naive implementation and my strategy uses 3 loop controls.