

HW2

b09902004 資工二 郭懷元

2.8

```
addi x30, x10, 8 → x30 = A+1
addi x31, x10, 0 → x31 = A
sd x31, 0(x30) → *(x30) = x31 → *(A+1) = A → A[1] = A
ld x30, 0(x30) → x30 = *(x30) → x30 = *(A+1) = A
add x5, x30, x31 → f = A + A
```

```
A[1] = A;
f = A[1] + A;
```

2.9

instruction	opcode	rs1	rd	rs2	imm	funct3	funct7
addi x30, x10, 8	010011	01010	11110		0x008	000	
addi x31, x10, 0	010011	01010	11111		0x000	000	
sd x31, 0(x30)	100011	11110		11111	0x000	011	
ld x30, 0(x30)	000011	11110	11110		0x000	011	
add x5, x30, x31	110011	11110	00110	11111		000	000

2.16

2.16.1

`funct7`, `funct3`, `opcode` : These bit fields might increase in size to accommodate the four times as many instructions.

`rs2`, `rs1`, `rd` : These bit fields should increase from 5 bits to 7 bits for the 128 registers.

2.16.2

`funct3`, `opcode` : These bit fields might increase in size to accommodate the four times as many instructions.

`rs1`, `rd` : These bit fields should increase from 5 bits to 7 bits for the 128 registers.

`imm` : This field doesn't need to change, because neither the number of registers or instructions have to do with `imm`.

2.16.3

Decrease in size: Because there are more registers and more instructions, some old instructions can now be combined into just a single instruction.

Increase in size: Because instructions now takes up more bits, for simple tasks that doesn't use many registers, the extra bits are wasted and take up unnecessary spaces.

Matrix Multiplication Report

Refs:

https://en.wikipedia.org/wiki/Loop_nest_optimization
<https://github.com/flame/how-to-optimize-gemm/wiki>

Result


Cycle count: 5630353

```
問題 輸出 终端機 偵錯主控台 連接埠

root@b06803fdad65:~/Problems/matrix# make run
riscv64-unknown-elf-gcc -O3 -o matrix matrix.c matrix.s
root@b06803fdad65:~/Problems/matrix# make test
spike pk ./matrix
bbl loader
Took 5630353 cycles
root@b06803fdad65:~/Problems/matrix#
```

Strategy

- Dot 4 rows of A and 4 columns of B at a time, so that when a value from A or B is read from memory to register, it can be used multiple times.



- Use pointers and indirect addressing to access A , B , and C .
- Only do the modulo operation before storing to C , because we are using `unsigned short`, and overflows in `unsigned` variables are essentially modulo.
- In the `for` loop of `k`, do two iterations at once, reducing call for branching operations.
- When inside the `for` loop of `i`, put `i` from register to stack, and use this register for as temporary register for calculation.
- I tried memory blocking to keep data in L1 and L2 cache, but somehow it didn't work, and the performance is even worse.
- Registers are barely enough. Unrolling some loops could avoid the use of `gp`, `tp`, and `ra`.

reg	usage	reg	usage
ra	k	s7	B[k][j+1]
t0	i / tmp	s8	B[k][j+2]
t1	j	s9	B[k][j+3]
t2	acc20	s10	acc31
t3	acc21	s11	acc32
t4	acc22	a0	A
t5	acc23	a1	B
t6	acc30	a2	C
s0	acc00	a3	acc33
s1	acc01	a4	acc10
s2	acc02	a5	acc11
s3	acc03	a6	acc12
s4	A[i][k]	a7	acc13
s5	A[i+1][k]	gp	A[i+2][k]
s6	B[k][j]	tp	A[i+3][k]